

Final_Project_4

March 10, 2023

Geographical and Time

```
[1]: import sys
      print(sys.version)
```

3.8.15 | packaged by conda-forge | (default, Nov 22 2022, 08:46:39)
[GCC 10.4.0]

```
[2]: spark.version
```

```
[2]: '3.1.3'
```

```
[3]: import pandas as pd
      import numpy as np
      pd.set_option('display.max_colwidth', None)
      pd.reset_option('display.max_rows')
      from itertools import compress
      from pyspark.sql.functions import *
      from pyspark.sql.types import *
      import seaborn as sns
      import matplotlib.pyplot as plt
      warnings.filterwarnings(action='ignore')
```

```
[4]: from pyspark.sql import SparkSession
      from pyspark import SparkContext
      from pyspark.sql import SQLContext
      from pyspark.sql import Row
      from pyspark.sql.functions import col
```

```
[5]: spark.conf.set("spark.sql.repl.eagerEval.enabled", True)
```

```
[6]: %%time
      twitter = spark.read.parquet('gs://chen26-bdp/original_data')
```

CPU times: user 6.35 ms, sys: 3.72 ms, total: 10.1 ms
Wall time: 8.38 s

23/03/05 02:19:11 WARN org.apache.spark.sql.catalyst.util.package: Truncated the string representation of a plan since it was too large. This behavior can be adjusted by setting 'spark.sql.debug.maxToStringFields'.

```
[7]: keywords = ['college', 'high', 'university', 'students',
                , 'public', 'private', 'secondary', 'primary', 'education',
                ↪ 'undergraduate', 'graduate']
#filter out rows that do not contain words in keywords
twitter = twitter.withColumn('lower', lower(col('text')))
filter_twitter = twitter.filter(col('lower').rlike('|'.join(keywords)))

twitter_eng = filter_twitter.filter(col('lang') == 'en')
from pyspark.sql import functions as F
from pyspark.sql import types as t
from pyspark.sql.types import ArrayType, IntegerType, BooleanType

eng_ord=F.udf(lambda x: [ord(a) for a in x],t.ArrayType(IntegerType()))

def english_filter(x):
    for index in range(len(x)):
        if x[index] > 128:
            return False
        else:
            return True

filter_udf = F.udf(english_filter, BooleanType())
tweets = twitter_eng.filter(filter_udf(eng_ord('text')) == True)
```

```
[8]: time = tweets.select('created_at')
```

```
[9]: time.show(10, truncate = False)
```

[Stage 1:>

(0 + 1) / 1]

```
+-----+
|created_at|
+-----+
|Sun Dec 18 00:53:14 +0000 2022|
|Sun Dec 18 00:53:15 +0000 2022|
|Sun Dec 18 00:53:16 +0000 2022|
|Sun Dec 18 00:53:17 +0000 2022|
|Sun Dec 18 00:53:18 +0000 2022|
|Sun Dec 18 00:53:18 +0000 2022|
|Sun Dec 18 00:53:19 +0000 2022|
|Sun Dec 18 00:53:19 +0000 2022|
|Sun Dec 18 00:53:22 +0000 2022|
|Sun Dec 18 00:53:22 +0000 2022|
+-----+
```

only showing top 10 rows

```
[10]: from pyspark.sql.functions import from_unixtime, date_format
spark.sql("set spark.sql.legacy.timeParserPolicy=LEGACY")

time_new = time.withColumn("timestamp_unix",
    ↪from_unixtime(unix_timestamp("created_at", "EEE MMM dd HH:mm:ss Z yyyy")))
time_new = time_new.withColumn("date", date_format("timestamp_unix",
    ↪"yyyy-MM-dd"))
```

```
[11]: # Show the updated DataFrame
time_new.show()
```

[Stage 2:>

(0 + 1) / 1]

```
+-----+-----+-----+
|      created_at|timestamp_unix|      date|
+-----+-----+-----+
|Sun Dec 18 00:53:...|2022-12-18 00:53:14|2022-12-18|
|Sun Dec 18 00:53:...|2022-12-18 00:53:15|2022-12-18|
|Sun Dec 18 00:53:...|2022-12-18 00:53:16|2022-12-18|
|Sun Dec 18 00:53:...|2022-12-18 00:53:17|2022-12-18|
|Sun Dec 18 00:53:...|2022-12-18 00:53:18|2022-12-18|
|Sun Dec 18 00:53:...|2022-12-18 00:53:18|2022-12-18|
|Sun Dec 18 00:53:...|2022-12-18 00:53:19|2022-12-18|
|Sun Dec 18 00:53:...|2022-12-18 00:53:19|2022-12-18|
|Sun Dec 18 00:53:...|2022-12-18 00:53:22|2022-12-18|
|Sun Dec 18 00:53:...|2022-12-18 00:53:22|2022-12-18|
|Sun Dec 18 00:53:...|2022-12-18 00:53:23|2022-12-18|
|Sun Dec 18 00:53:...|2022-12-18 00:53:27|2022-12-18|
|Sun Dec 18 00:53:...|2022-12-18 00:53:28|2022-12-18|
|Sun Dec 18 00:53:...|2022-12-18 00:53:28|2022-12-18|
|Sun Dec 18 00:53:...|2022-12-18 00:53:29|2022-12-18|
|Sun Dec 18 00:53:...|2022-12-18 00:53:29|2022-12-18|
|Sun Dec 18 00:53:...|2022-12-18 00:53:31|2022-12-18|
|Sun Dec 18 00:53:...|2022-12-18 00:53:32|2022-12-18|
|Sun Dec 18 00:53:...|2022-12-18 00:53:33|2022-12-18|
|Sun Dec 18 00:53:...|2022-12-18 00:53:33|2022-12-18|
+-----+-----+-----+
```

only showing top 20 rows

```
[12]: tweets_day = tweets.select('id', 'created_at')
```

```

tweets_day = tweets_day.withColumn("timestamp_unix",
    ↪from_unixtime(unix_timestamp("created_at", "EEE MMM dd HH:mm:ss Z yyyy")))
tweets_day = tweets_day.withColumn("date", date_format("timestamp_unix",
    ↪"yyyy-MM-dd"))
tweets_by_day = tweets_day.drop('created_at', 'timestamp_unix')

```

```

[13]: # Show the updated DataFrame
      tweets_by_day.show()

```

[Stage 3:>

(0 + 1) / 1]

```

+-----+-----+
|          id|      date|
+-----+-----+
|1604278583291256832|2022-12-18|
|1604278587183415297|2022-12-18|
|1604278590614315008|2022-12-18|
|1604278595215671296|2022-12-18|
|1604278597019115520|2022-12-18|
|1604278598155702273|2022-12-18|
|1604278600655507456|2022-12-18|
|1604278604354834433|2022-12-18|
|1604278614618279936|2022-12-18|
|1604278614748528642|2022-12-18|
|1604278620993757186|2022-12-18|
|1604278633853583362|2022-12-18|
|1604278640241278977|2022-12-18|
|1604278641076092931|2022-12-18|
|1604278643693424643|2022-12-18|
|1604278645501181952|2022-12-18|
|1604278651373199361|2022-12-18|
|1604278654892183553|2022-12-18|
|1604278659983937536|2022-12-18|
|1604278661267574785|2022-12-18|
+-----+-----+

```

only showing top 20 rows

```

[14]: total_by_day = tweets_by_day.groupby('date').agg(count('*').
    ↪alias('daily_tweets'))
      day_total = total_by_day.orderBy('date', ascending = True)
      day_total.show()

```

[Stage 6:=====>

(7 + 8) / 15]

```

+-----+-----+
|      date|daily_tweets|

```

```

+-----+-----+
|2022-04-05|    109939|
|2022-04-06|    141312|
|2022-04-07|    141789|
|2022-04-08|    137075|
|2022-04-09|    109139|
|2022-04-10|     98880|
|2022-04-11|    142839|
|2022-04-12|    137345|
|2022-04-13|    138168|
|2022-04-14|    118257|
|2022-04-15|    106518|
|2022-04-16|    100385|
|2022-04-17|     84230|
|2022-04-18|    115006|
|2022-04-19|    135267|
|2022-04-20|    126609|
|2022-04-21|    124602|
|2022-04-22|    121740|
|2022-04-23|     97615|
|2022-04-24|    101238|
+-----+-----+
only showing top 20 rows

```

```

[16]: %%time
      day_total.write.format("parquet").\
      mode('overwrite').\
      save('gs://chen26-bdp/day_tweets')

```

CPU times: user 892 ms, sys: 213 ms, total: 1.1 s
 Wall time: 6min 23s

```

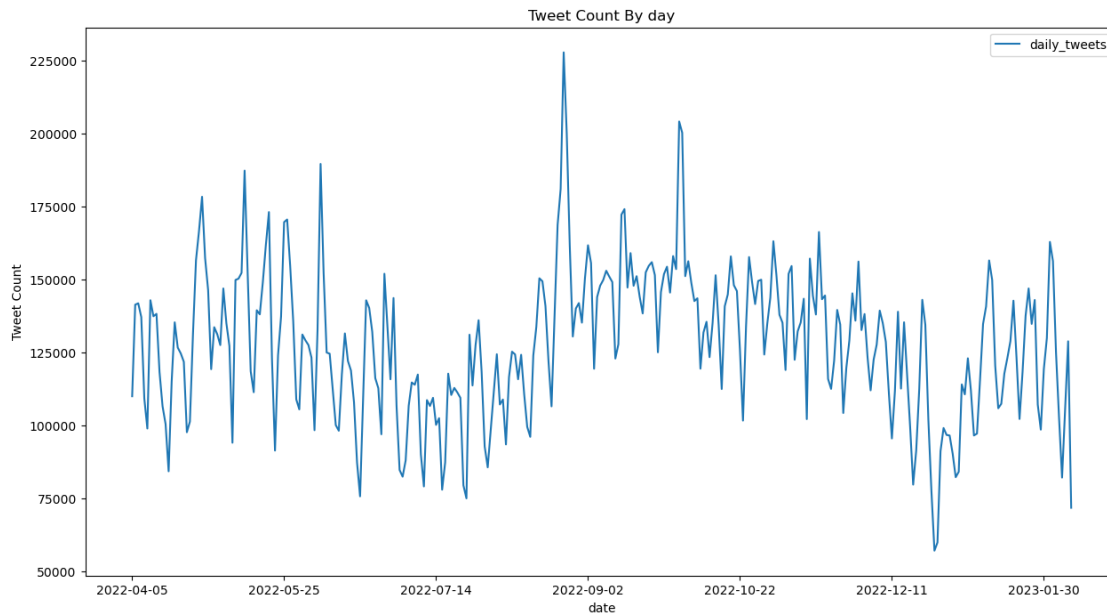
[15]: day = day_total.toPandas()

```

```

[17]: day.plot(kind="line",x= 'date', y="daily_tweets", figsize=(15, 8))
      plt.ylabel("Tweet Count")
      plt.title('Tweet Count By day')
      plt.show()

```



```
[18]: day
```

```
[18]:
```

	date	daily_tweets
0	2022-04-05	109939
1	2022-04-06	141312
2	2022-04-07	141789
3	2022-04-08	137075
4	2022-04-09	109139
..
305	2023-02-04	102690
306	2023-02-05	82069
307	2023-02-06	105679
308	2023-02-07	128770
309	2023-02-08	71723

```
[310 rows x 2 columns]
```

```
[19]: month_total = day_total.withColumn('month', month(col('date')))
```

```
[ ]: month_total.show()
```

```
[Stage 25:=====>
```

```
(5 + 8) / 13]
```

```
+-----+-----+-----+
|      date|daily_tweets|month|
+-----+-----+-----+
|2022-04-05|      109939|    4|
|2022-04-06|      141312|    4|
```

2022-04-07	141789	4
2022-04-08	137075	4
2022-04-09	109139	4
2022-04-10	98880	4
2022-04-11	142839	4
2022-04-12	137345	4
2022-04-13	138168	4
2022-04-14	118257	4
2022-04-15	106518	4
2022-04-16	100385	4
2022-04-17	84230	4
2022-04-18	115006	4
2022-04-19	135267	4
2022-04-20	126609	4
2022-04-21	124602	4
2022-04-22	121740	4
2022-04-23	97615	4
2022-04-24	101238	4

+-----+-----+-----+

only showing top 20 rows

```
[ ]: %%time
month_total.write.format("parquet").\
mode('overwrite').\
save('gs://chen26-bdp/month_tweets')
```

```
[7]: month_total = spark.read.parquet('gs://chen26-bdp/month_tweets')
```

```
[9]: month_total_t = month_total.groupby('month').agg(sum('daily_tweets').
↳alias('monthly_tweets'))
```

```
[10]: month = month_total_t.toPandas()
```

```
[15]: month = month.sort_values(by = 'month')
```

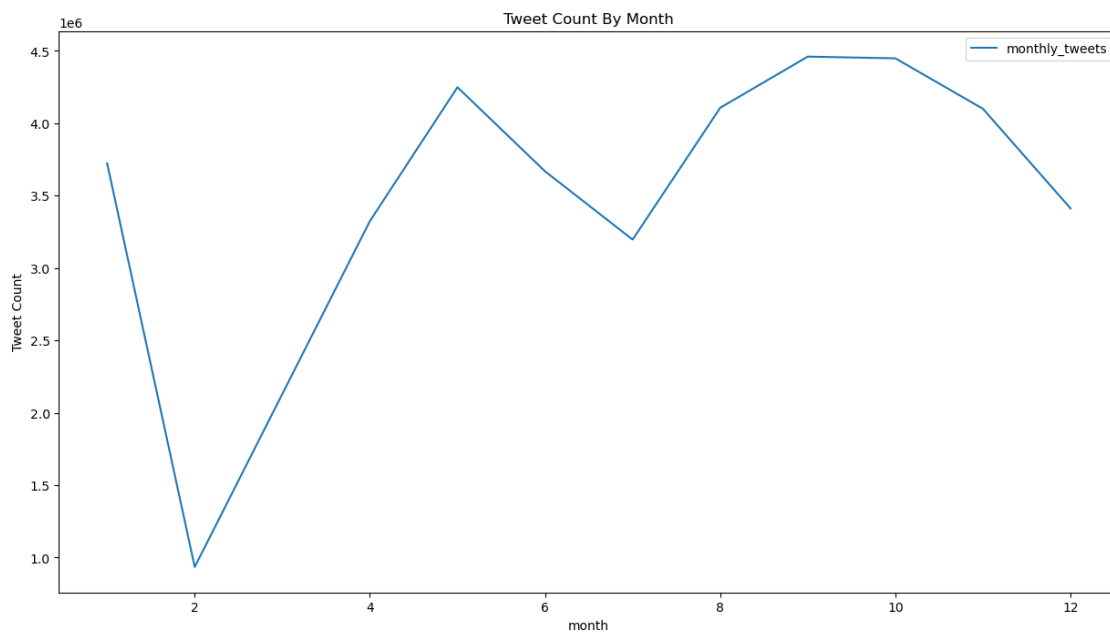
```
[16]: month
```

```
[16]:
```

	month	monthly_tweets
1	1	3722912
9	2	935033
5	4	3323905

3	5	4248605
2	6	3667682
6	7	3196325
10	8	4106759
4	9	4460274
7	10	4448331
8	11	4099704
0	12	3411590

```
[17]: month.plot(kind="line",x= 'month', y="monthly_tweets", figsize=(15, 8))
plt.ylabel("Tweet Count")
plt.title('Tweet Count By Month')
plt.show()
```



```
[29]: spark.sql("set spark.sql.legacy.timeParserPolicy=LEGACY")
tweets_week = tweets.select('id', 'created_at')
tweets_week = tweets_week.withColumn("timestamp_unix",
    →from_unixtime(unix_timestamp("created_at", "EEE MMM dd HH:mm:ss Z yyyy")))
tweets_week = tweets_week.withColumn('day_of_week',
    →date_format(to_date('created_at', 'EEE MMM dd HH:mm:ss Z yyyy'), 'EEEE'))
#tweets_by_day = tweets_day.drop('created_at', 'timestamp_unix')
```

```
[31]: tweets_week = tweets_week.drop('timestamp_unix')
```

```
[32]: day_of_week = tweets_week.groupby('day_of_week').agg(count('id'))
```



```
[ ]: dw = day_of_week.toPandas()
```

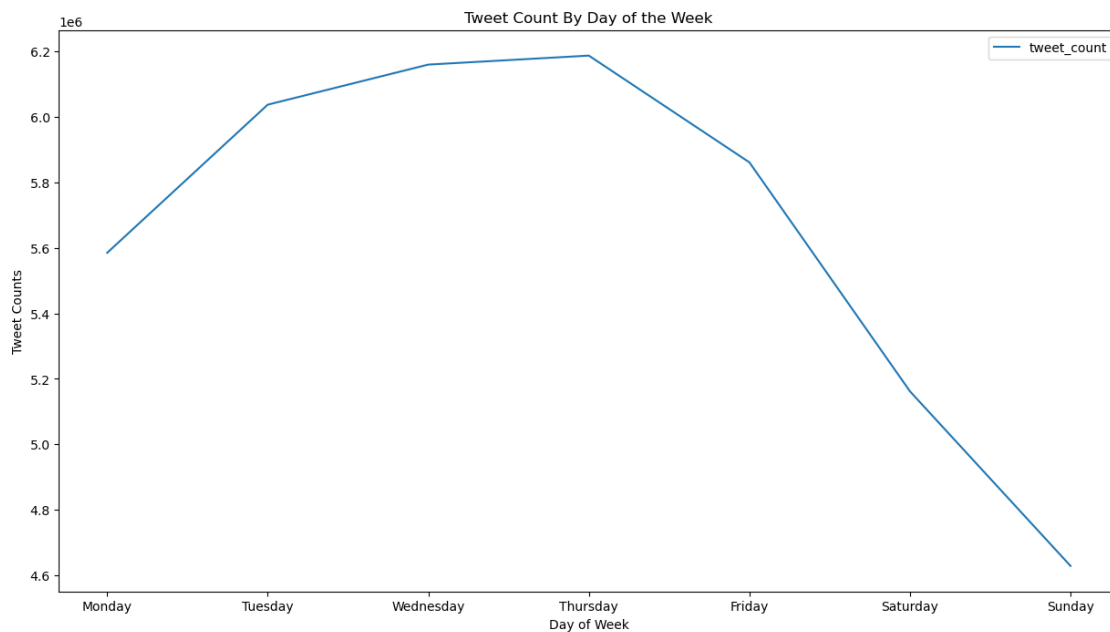
```
[35]: dw.sort_values(by = 'day_of_week')
```

```
[35]:  day_of_week  count(id)
      2    Friday    5861154
      5    Monday    5584926
      4   Saturday    5162074
      6    Sunday    4629016
      3  Thursday    6187074
      1   Tuesday    6037232
      0  Wednesday    6159644
```

```
[37]: days = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']
      count = [5584926, 6037232, 6159644, 6187074, 5861154, 5162074, 4629016]

      tweet_week = pd.DataFrame({'Days': days, 'tweet_count': count})
```

```
[39]: tweet_week.plot(kind="line", x= 'Days', y="tweet_count", figsize=(15, 8))
      plt.ylabel("Tweet Counts")
      plt.xlabel('Day of Week')
      plt.title('Tweet Count By Day of the Week')
      plt.show()
```



Geographic

```
[ ]: # tweets location

[40]: tweets_loc = tweets.select('id', tweets.place.country.alias('tweet_country'),  
    ↪ tweets.place.full_name.alias('tweet_location'))

[42]: tweet_country = tweets_loc.select('id', 'tweet_country')

[51]: country = tweet_country.filter(col('tweet_country').isNotNull())

[59]: t_by_country = country.groupby(col('tweet_country')).agg(count('*').  
    ↪ alias('num_tweets'))


[61]: top50_country = t_by_country.orderBy(col('num_tweets').desc()).limit(50)

[63]: top50_country_df = top50_country.toPandas()

[72]: ax = top50_country_df.plot(x='tweet_country',y='num_tweets', kind='barh',  
    ↪ figsize = (10, 8))

    ax.set_ylabel('Country')
    ax.set_xlabel('Total tweets')
    ax.set_title('Top 50 Country of Total Number of Tweets')
    ax.invert_yaxis()

    # show the plot
    plt.show()
```



```
[10]: tweet_location = tweets_loc.select('id', 'tweet_location')
```

```
[11]: location = tweet_location.filter(col('tweet_location').isNotNull())
```

```
[16]: location_c = location.withColumn("loc", split(location.tweet_location, ", ").
    ↳ getItem(0))
```

```
[17]: location_c
```

```
[17]: +-----+-----+-----+
|      id| tweet_location|      loc|
+-----+-----+-----+
|1555124473338253315| Punjab, Pakistan| Punjab|
|1555124856861360128| Bhubaneshwar, India| Bhubaneshwar|
|1555125171669061632| India| India|
|1555125396970115072| Greengairs, Scotland| Greengairs|
|1530714698655596545| Mount Vernon Pres...| Mount Vernon Pres...|
|1530714877664403457| Marshfield, MA| Marshfield|
|1530715122515296256| Little Falls, NJ| Little Falls|
|1530715510123466752| Sarasota, FL| Sarasota|
|1530715614649647104| Georgia, USA| Georgia|
|1530716000383025159| Virginia Beach, VA| Virginia Beach|
|1530716062416875521| Lake Central High...| Lake Central High...|
```

1533456562458464261	Philadelphia, PA	Philadelphia
1533456594264104961	North Carolina, USA	North Carolina
1533456664585719809	North Olmsted, OH	North Olmsted
1533456934665211904	Collinsville, IL	Collinsville
1533457154220470274	Philadelphia, PA	Philadelphia
1530540890221158402	Maine, USA	Maine
1530541106454200320	Nsukka, Nigeria	Nsukka
1530541114918313984	Legacy Park	Legacy Park
1530541231679356928	Kentucky, USA	Kentucky

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+

only showing top 20 rows

```
[19]: t_by_location = location_c.groupby(col('loc')).agg(count('*').
      ↪alias('num_tweets'))
```

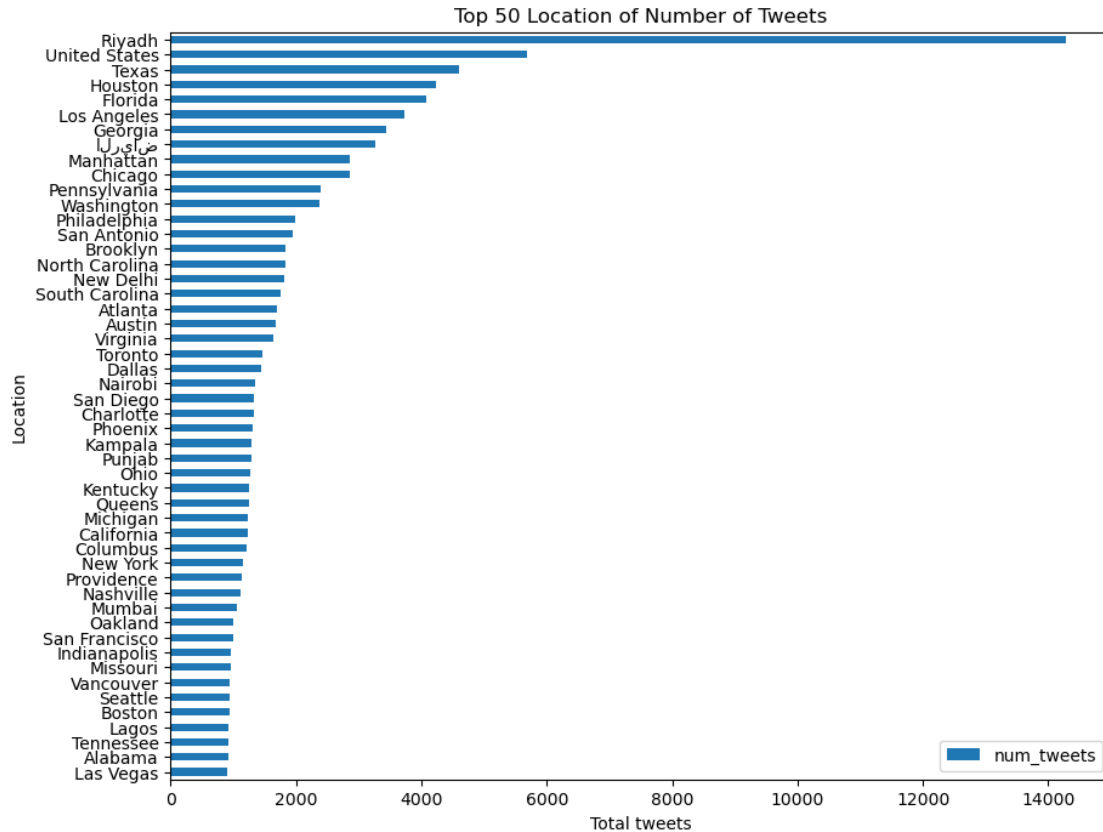
```
[20]: top50_loc = t_by_location.orderBy(col('num_tweets').desc()).limit(50)
```

```
[22]: top50_loc_df = top50_loc.toPandas()
```

```
[24]: ax = top50_loc_df.plot(x='loc',y='num_tweets', kind='barh', figsize = (10, 8))

ax.set_ylabel('Location')
ax.set_xlabel('Total tweets')
ax.set_title('Top 50 Location of Number of Tweets')
ax.invert_yaxis()

# show the plot
plt.show()
```



```
[25]: # remove locations with only country
filt_location = location.filter(col("tweet_location").like("%,%"))
```

```
[27]: location_f = filt_location.withColumn("loc", split(locationtweet_location, ",,
↪").getItem(0))
```

```
[29]: location_f
```

```
[29]: +-----+-----+-----+
|          id|    tweet_location|          loc|
+-----+-----+-----+
|1555124473338253315|    Punjab, Pakistan|    Punjab|
|1555124856861360128| Bhubaneshwar, India| Bhubaneshwar|
|1555125396970115072| Greengairs, Scotland| Greengairs|
|1530714877664403457|    Marshfield, MA|    Marshfield|
|1530715122515296256|    Little Falls, NJ|    Little Falls|
|1530715510123466752|    Sarasota, FL|    Sarasota|
|1530715614649647104|    Georgia, USA|    Georgia|
|1530716000383025159| Virginia Beach, VA| Virginia Beach|
|1533456562458464261| Philadelphia, PA| Philadelphia|
```

1533456594264104961	North Carolina, USA	North Carolina
1533456664585719809	North Olmsted, OH	North Olmsted
1533456934665211904	Collinsville, IL	Collinsville
1533457154220470274	Philadelphia, PA	Philadelphia
1530540890221158402	Maine, USA	Maine
1530541106454200320	Nsukka, Nigeria	Nsukka
1530541231679356928	Kentucky, USA	Kentucky
1530541346032861184	Nawabshah, Pakistan	Nawabshah
1530541387443122176	Chandigarh, India	Chandigarh
1566012216994070528	Kota, India	Kota
1566012298983923713	Maharashtra, India	Maharashtra

+-----+-----+-----+-----+

only showing top 20 rows

```
[30]: t_by_location_f = location_f.groupby(col('loc')).agg(count('*').
      ↪alias('num_tweets'))
```

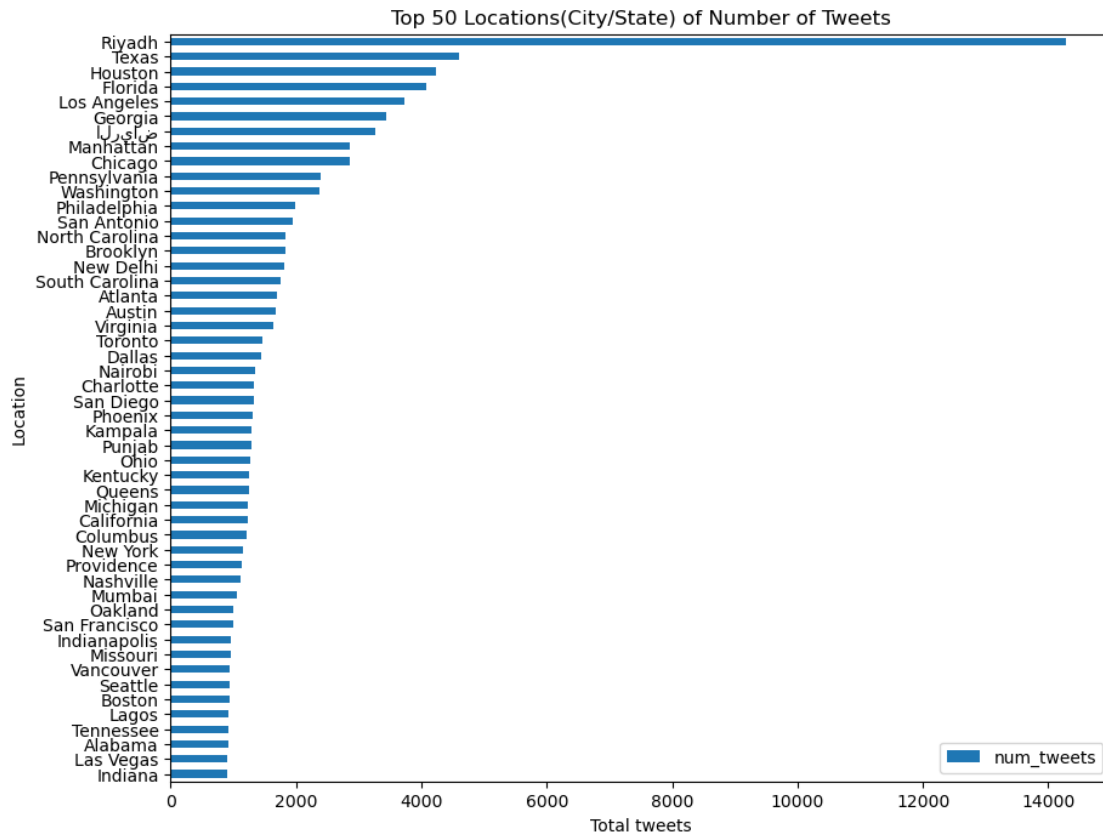
```
[31]: top50_loc_f = t_by_location_f.orderBy(col('num_tweets').desc()).limit(50)
```

```
[32]: top50_loc_f_df = top50_loc_f.toPandas()
```

```
[81]: ax = top50_loc_f_df.plot(x='loc',y='num_tweets', kind='barh', figsize = (10,
      ↪8))

ax.set_ylabel('Location')
ax.set_xlabel('Total tweets')
ax.set_title('Top 50 Locations(City/State) of Number of Tweets')
ax.invert_yaxis()

# show the plot
plt.show()
```



```
[ ]: # check why many tweets come from riyadh
```

```
[43]: riyadh = tweets.filter(tweets.place.full_name == 'Riyadh, Kingdom of Saudi Arabia')
```

```
[51]: spark.sql("set spark.sql.legacy.timeParserPolicy=LEGACY")
      riyadh_timeline = riyadh.select('id', 'created_at')
      riyadh_month = riyadh_timeline.withColumn("timestamp_unix",
        from_unixtime(unix_timestamp("created_at", "EEE MMM dd HH:mm:ss Z yyyy")))
      riyadh_month = riyadh_month.withColumn("date", date_format("timestamp_unix",
        "yyyy-MM-dd"))
      riyadh_month = riyadh_month.withColumn('month', month('date'))
```

```
[56]: r_d = riyadh_month.select('id', 'date')
      r_tweets = r_d.groupby('date').agg(count('*'))
```

```
[ ]: r_tweets_day = r_tweets.orderBy('date', ascending = True)
      r_tweets_day.show()
```

[Stage 79:=====>

(9 + 4) / 13]

date	count(1)
2022-04-05	13
2022-04-06	8
2022-04-07	7
2022-04-08	5
2022-04-09	7
2022-04-10	5
2022-04-11	16
2022-04-12	36
2022-04-13	13
2022-04-14	11
2022-04-15	7
2022-04-16	7
2022-04-17	16
2022-04-18	23
2022-04-19	40
2022-04-20	40
2022-04-21	41
2022-04-22	27
2022-04-23	9
2022-04-24	8

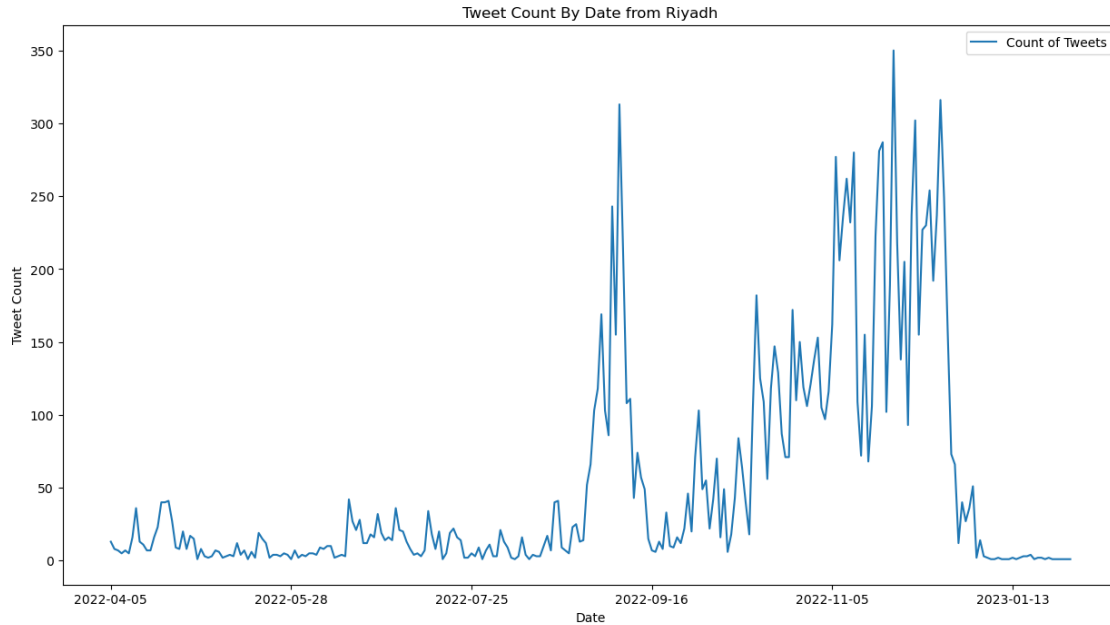
only showing top 20 rows

```
[59]: rt_df = r_tweets_day.toPandas()
```

```
[63]: ax = rt_df.plot(kind="line", x='date', y="count(1)", figsize=(15, 8))
ax.set_ylabel("Tweet Count")
ax.set_xlabel('Date')
ax.set_title('Tweet Count By Date from Riyadh')

# Adding Legend and Renaming
ax.legend(['Count of Tweets'])

plt.show()
```

```
[ ]: # peak 1: Jul, Aug: Biden saudi visit
    # peak 2: Dec: world cup
```

```
[ ]:
```

```
[ ]:
```

```
[ ]: # twitter user location
```

```
[64]: user_loc = tweets.select('id', tweets.user.location.alias('user_location'),
    ↳ tweets.user.id.alias('user_id'))
```

```
[66]: u_loc = user_loc.filter(col('user_location').isNotNull())
```

```
[68]: u_loc_count = u_loc.groupby('user_location').agg(count('user_id'))
```

```
[71]: lc_desc = u_loc_count.orderBy(col('count(user_id)').desc())
```

```
[72]: lc = lc_desc.withColumn("loction", split(lc_desc.user_location, ", ").
    ↳ getItem(0))
```

```
[74]: lc_f = lc.filter(col("user_location").like("%,%"))
```

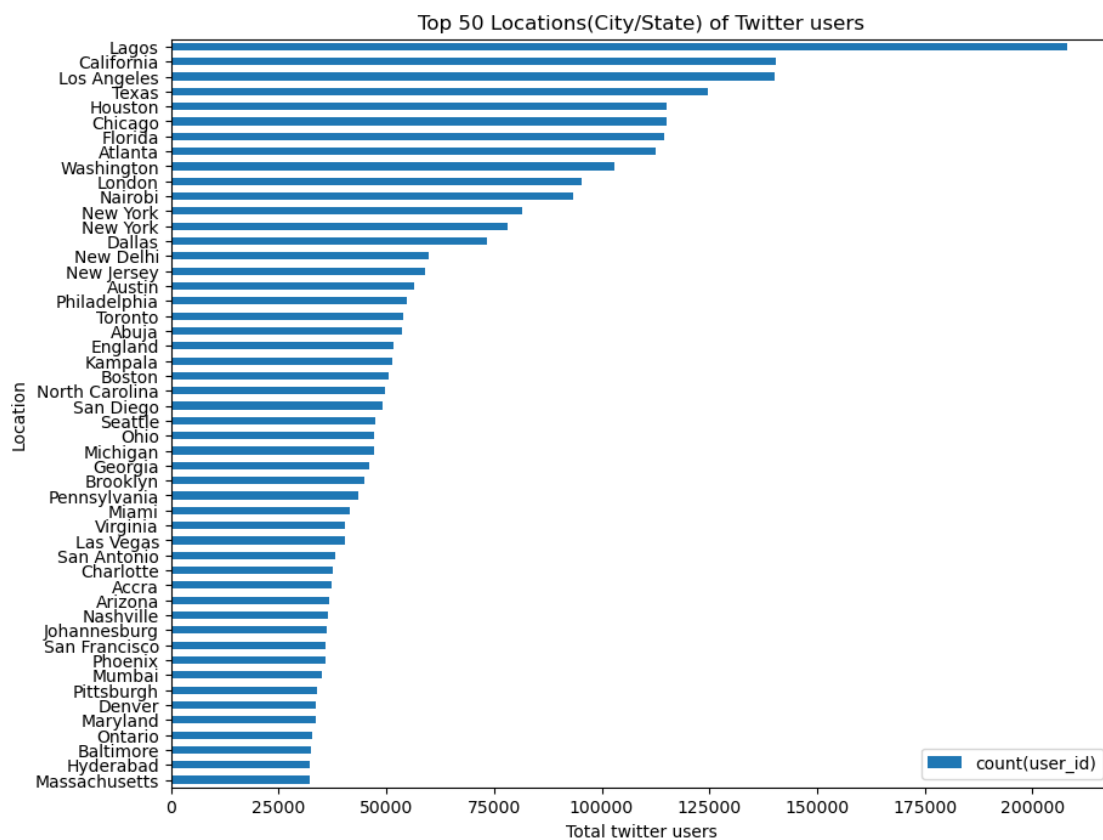
```
[77]: l_df = lc_f.limit(50).toPandas()
    ax = l_df.plot(x='loction', y='count(user_id)', kind='barh', figsize=(10, 8))
```

```

ax.set_ylabel('Location')
ax.set_xlabel('Total twitter users')
ax.set_title('Top 50 Locations(City/State) of Twitter users')
ax.invert_yaxis()

# show the plot
plt.show()

```



```
[ ]: # focus peaks in trend: may
```

```

[9]: spark.sql("set spark.sql.legacy.timeParserPolicy=LEGACY")
    timeline = tweets.select('id', 'created_at', tweets.place.full_name.
        ↳alias('tweet_location'))
    timeline = timeline.withColumn("timestamp_unix",
        ↳from_unixtime(unix_timestamp("created_at", "EEE MMM dd HH:mm:ss Z yyyy")))
    timeline = timeline.withColumn("date", date_format("timestamp_unix",
        ↳"yyyy-MM-dd"))
    timeline = timeline.withColumn('month', month('date'))

```

```
[11]: timeline_loc = timeline.filter(col('tweet_location').isNotNull())
```

```
[13]: timeline_loc_may = timeline_loc.filter(col('month') == 5).select('id',
    ↳ 'tweet_location', 'date')

[15]: timeline_l_may = timeline_loc_may.withColumn('loc', split(timeline_loc_may.
    ↳ tweet_location, ", ").getItem(0))

[20]: may = timeline_l_may.groupby('loc').agg(count('id').alias('num_tweets')).
    ↳ orderBy(col('num_tweets').desc()).limit(50)

[ ]: may.show()
```

[Stage 19:=====> (7 + 10) / 17]

```
+-----+-----+
|          loc|num_tweets|
+-----+-----+
| United States|      4041|
|      Houston|      443|
|      Georgia|      410|
| Los Angeles|      370|
|      Florida|      354|
|      Texas|      339|
|      Manhattan|      303|
|      Chicago|      278|
| Washington|      241|
| Philadelphia|      209|
| Pennsylvania|      199|
|      Brooklyn|      192|
|North Carolina|      187|
| San Antonio|      173|
|      Virginia|      171|
|      Atlanta|      170|
|      Austin|      168|
|South Carolina|      154|
|      New Delhi|      153|
|      Riyadh|      147|
+-----+-----+
only showing top 20 rows
```

Topics

```
[ ]: # "Florida math book ban"

[8]: dt = tweets.select([tweets.created_at,
    tweets.id_str.alias('tweet_id'),
    tweets.user['id'].alias('user_id'),
```

```

    tweets.user['screen_name'].alias('user_name'),
    tweets.user['verified'].alias('verified'),
    tweets.user['followers_count'].alias('followers_count'),
    tweets.user['description'].alias('user_description'),
    tweets.user['location'].alias('user_location'),
    tweets.text,
    tweets.retweeted_status.retweet_count.alias('rt_count'),
    tweets.retweeted_status.favorite_count.alias('rt_fav'),
    tweets.retweeted_status.quote_count.alias('rt_quo'),
    tweets.retweeted_status.retweeted.alias('retweet'),
    tweets.retweeted_status.user['name'].alias('rt_user_name'))]]

```

```
[ ]: # filter texts to find those that contain "Florida math book ban"
```

```

[15]: fl_key = ['Florida', 'math', 'book', 'ban']
      #filter out rows that do not contain words in keywords
      dt = dt.withColumn('lower', lower(col('text')))
      fl_dt = dt.filter(col('lower').rlike('|'.join(fl_key)))

```

```
[16]: fl_dt.count()
```

```
[16]: 1620798
```

```
[27]: fl_dt_loc = fl_dt.filter(col('user_location').isNotNull())
```

```

[28]: spark.sql("set spark.sql.legacy.timeParserPolicy=LEGACY")
      fl_may = fl_dt_loc.select('tweet_id', 'created_at', 'user_location')
      fl_may = fl_may.withColumn("timestamp_unix",
      ↪from_unixtime(unix_timestamp("created_at", "EEE MMM dd HH:mm:ss Z yyyy")))
      fl_may = fl_may.withColumn("date", date_format("timestamp_unix", "yyyy-MM-dd"))
      fl_may = fl_may.withColumn('month', month('date'))

```

```
[29]: fl_may
```

```

[29]: +-----+-----+-----+-----+
      --+-----+-----+
      |          tweet_id|          created_at|          user_location|
      timestamp_unix|          date|month|
      +-----+-----+-----+-----+
      --+-----+-----+
      |1604278645501181952|Sun Dec 18 00:53:...|    Capital city, LA|2022-12-18
      00:53:29|2022-12-18|    12|
      |1604278670776041472|Sun Dec 18 00:53:...|    United States|2022-12-18

```

```

00:53:35|2022-12-18|    12|
|1604278780037660674|Sun Dec 18 00:54:...|    Missouri, USA|2022-12-18
00:54:01|2022-12-18|    12|
|1604279313477632003|Sun Dec 18 00:56:...|    San Antonio, TX|2022-12-18
00:56:09|2022-12-18|    12|
|1604279364690104322|Sun Dec 18 00:56:...|Duluth,Minnesota USA|2022-12-18
00:56:21|2022-12-18|    12|
|1604279474551259136|Sun Dec 18 00:56:...|Stellenbosch, Sou...|2022-12-18
00:56:47|2022-12-18|    12|
|1604279577362178048|Sun Dec 18 00:57:...|    thy phrontistery|2022-12-18
00:57:11|2022-12-18|    12|
|1604279665757409280|Sun Dec 18 00:57:...| western Washington |2022-12-18
00:57:33|2022-12-18|    12|
|1604280491238096897|Sun Dec 18 01:00:...|    Nigeria|2022-12-18
01:00:49|2022-12-18|    12|
|1604280526268923906|Sun Dec 18 01:00:...|    Floyds Knobs, IN|2022-12-18
01:00:58|2022-12-18|    12|
|1604280555100684291|Sun Dec 18 01:01:...|    Atlanta, GA|2022-12-18
01:01:05|2022-12-18|    12|
|1604280601548394496|Sun Dec 18 01:01:...|    Myrtle Beach, SC|2022-12-18
01:01:16|2022-12-18|    12|
|1604280711128797184|Sun Dec 18 01:01:...|    Myrtle Beach, SC|2022-12-18
01:01:42|2022-12-18|    12|
|1604280809908846592|Sun Dec 18 01:02:...|England-Turkey- C...|2022-12-18
01:02:05|2022-12-18|    12|
|1610637225871511552|Wed Jan 04 14:00:...|An 804 girl livin...|2023-01-04
14:00:13|2023-01-04|    1|
|1610637264421355526|Wed Jan 04 14:00:...|    Myrtle Beach, SC|2023-01-04
14:00:22|2023-01-04|    1|
|1610637278346354689|Wed Jan 04 14:00:...|    United Kingdom|2023-01-04
14:00:26|2023-01-04|    1|
|1610637463474651137|Wed Jan 04 14:01:...|    Myrtle Beach, SC|2023-01-04
14:01:10|2023-01-04|    1|
|1610637510492803077|Wed Jan 04 14:01:...|    Firenze|2023-01-04
14:01:21|2023-01-04|    1|
|1610637652121845761|Wed Jan 04 14:01:...|    Midlands, England|2023-01-04
14:01:55|2023-01-04|    1|
+-----+-----+-----+-----+
--+-----+-----+
only showing top 20 rows

```

```
[30]: fl_m_count = fl_may.groupby('date').agg(count('*').alias('daily_tweets'))
```

```
[ ]: fl_m_count
```

```
[ ]: +-----+-----+
|      date|daily_tweets|
+-----+-----+
|2022-07-30|      2892|
|2022-09-03|      2933|
|2022-12-14|      2134|
|2022-05-04|      2538|
|2022-11-04|      6096|
|2022-11-27|      2718|
|2022-09-07|      7552|
|2022-05-16|      2270|
|2022-07-23|      2107|
|2022-12-07|      2906|
|2022-05-17|      3294|
|2022-10-13|      3035|
|2022-11-09|      2549|
|2022-10-24|      2952|
|2022-11-23|      2370|
|2022-08-17|      5485|
|2023-01-05|      4164|
|2022-07-04|      1992|
|2022-07-25|      2327|
|2022-06-03|      3175|
+-----+-----+
only showing top 20 rows
```

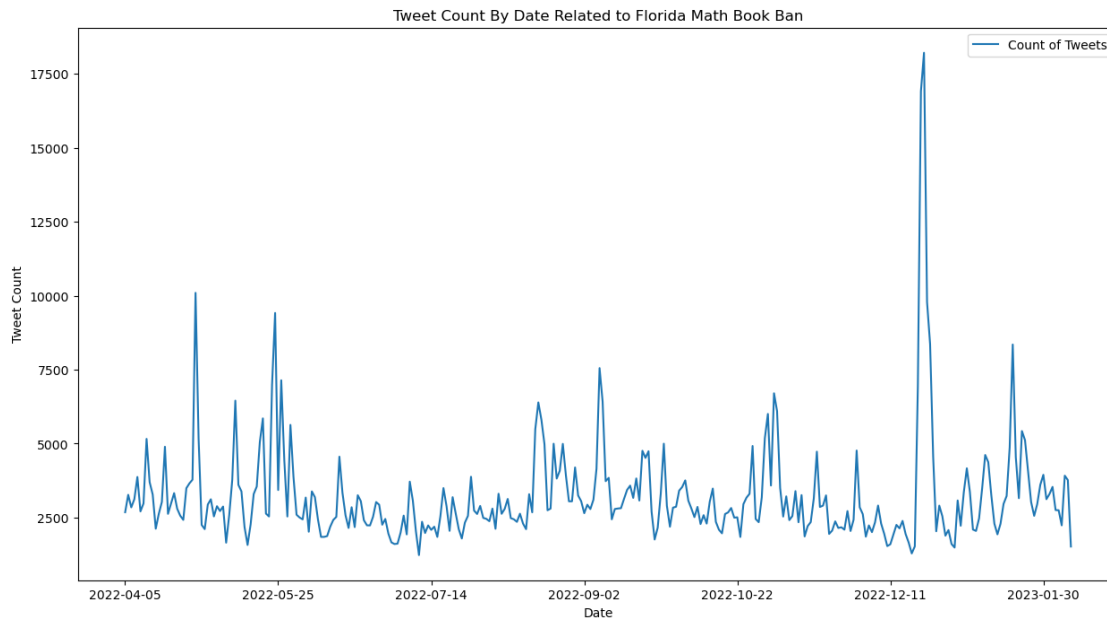
```
[32]: fl_count = fl_m_count.orderBy('date', ascending = True)
```

```
[33]: fl_df = fl_count.toPandas()
```

```
[36]: ax = fl_df.plot(kind="line", x='date', y="daily_tweets", figsize=(15, 8))
ax.set_ylabel("Tweet Count")
ax.set_xlabel('Date')
ax.set_title('Tweet Count By Date Related to Florida Math Book Ban')

# Adding Legend and Renaming
ax.legend(['Count of Tweets'])

plt.show()
```



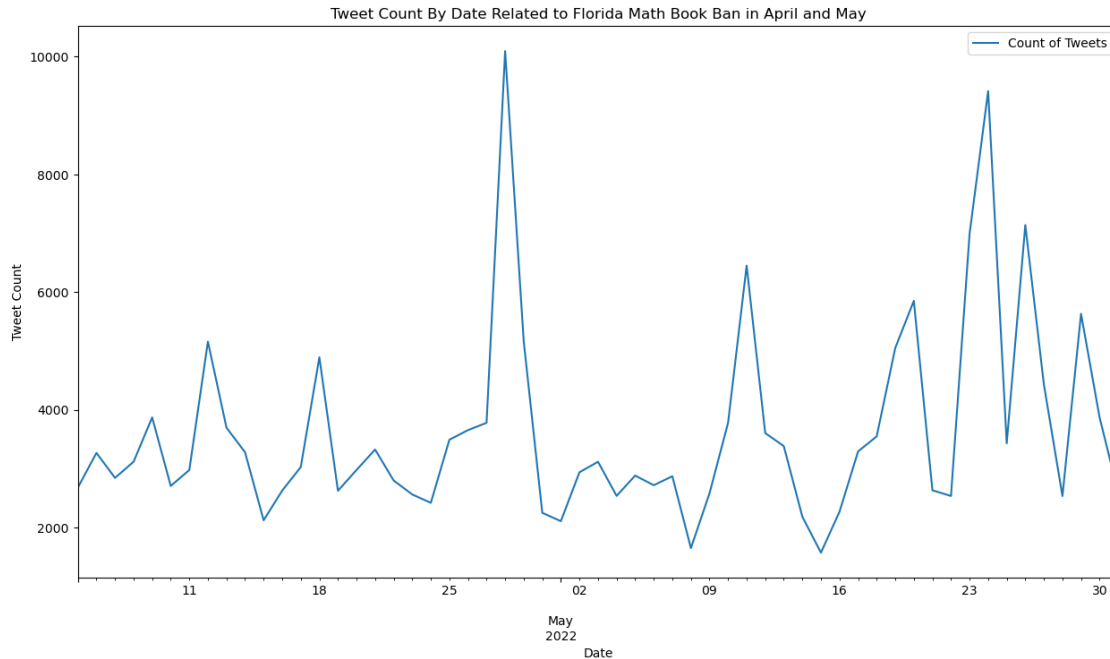
```
[48]: fl_df['date'] = pd.to_datetime(fl_df['date'])
      fl_df['M'] = fl_df['date'].dt.month
```

```
[55]: fl_df_4_5 = fl_df[(fl_df['M'] == 4) | (fl_df['M'] == 5)]
```

```
[57]: ax = fl_df_4_5.plot(kind="line", x='date', y="daily_tweets", figsize=(15, 8))
      ax.set_ylabel("Tweet Count")
      ax.set_xlabel('Date')
      ax.set_title('Tweet Count By Date Related to Florida Math Book Ban in April and
      ↪May')

      # Adding Legend and Renaming
      ax.legend(['Count of Tweets'])

      plt.show()
```



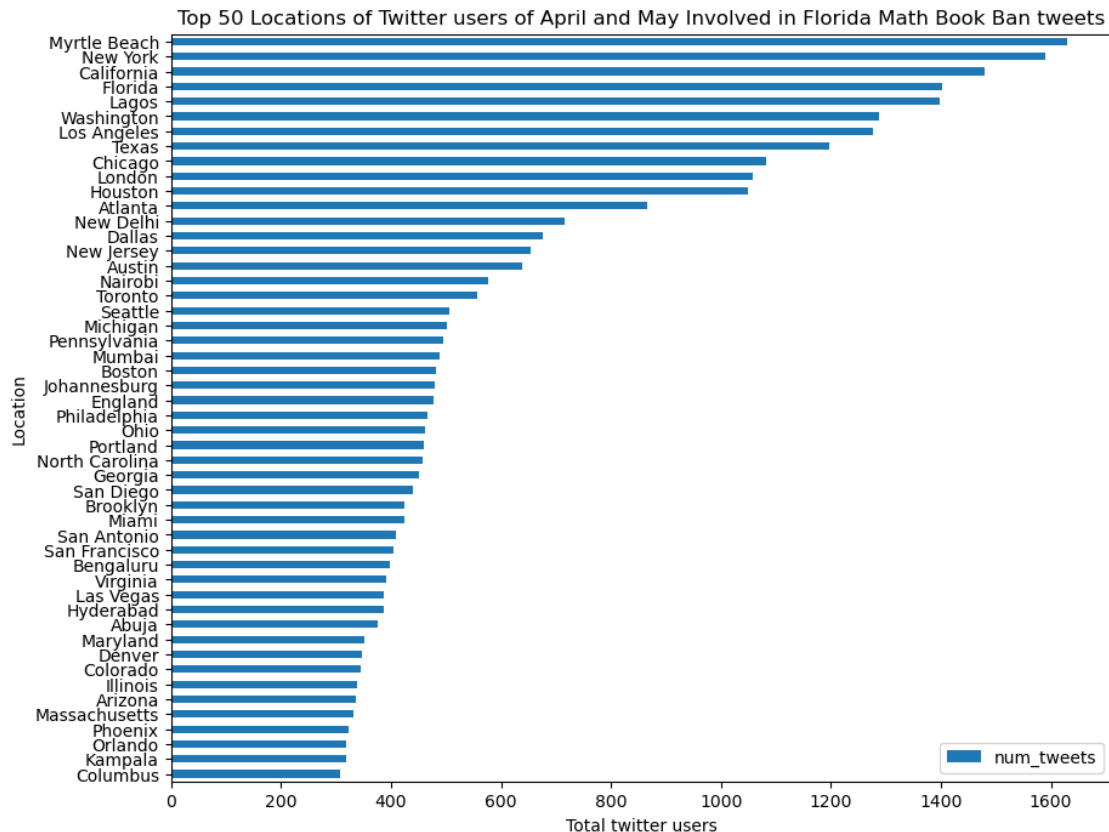
```
[58]: fl_4_5 = fl_may.filter(col('month').isin([4, 5]))
fl_4_5 = fl_4_5.withColumn('loc', split(fl_4_5.user_location, ", ").getItem(0))
fl_4_5 = fl_4_5.filter(col("user_location").like("%,%"))
fl_count = fl_4_5.groupby('loc').agg(count('*').alias('num_tweets'))
fl_desc = fl_count.orderBy(col('num_tweets').desc()).limit(50)
```

```
[59]: fl_desc_df = fl_desc.toPandas()
```

```
[70]: ax = fl_desc_df.plot(x='loc',y='num_tweets', kind='barh', figsize = (10, 8))

ax.set_ylabel('Location')
ax.set_xlabel('Total twitter users')
ax.set_title('Top 50 Locations of Twitter users of April and May Involved in_
↳Florida Math Book Ban tweets')
ax.invert_yaxis()

# show the plot
plt.show()
```

```
[ ]: spark.sql("set spark.sql.legacy.timeParserPolicy=LEGACY")
tdt = dt.filter(col('user_location').isNotNull())
tdt = tdt.select('tweet_id', 'created_at', 'user_location')
tdt = tdt.withColumn("timestamp_unix",
    →from_unixtime(unix_timestamp("created_at", "EEE MMM dd HH:mm:ss Z yyyy")))
tdt = tdt.withColumn("date", date_format("timestamp_unix", "yyyy-MM-dd"))
tdt = tdt.withColumn('month', month('date'))
tdt_4_5 = tdt.filter(col('month').isin([4, 5]))
tdt_4_5 = tdt_4_5.withColumn('loc', split(tdt_4_5.user_location, ", ").
    →getItem(0))
tdt_4_5 = tdt_4_5.filter(col("user_location").like("%,%"))
tdt_count = tdt_4_5.groupby('loc').agg(count('*').alias('num_tweets'))
tdt_desc = tdt_count.orderBy(col('num_tweets').desc()).limit(50)
tdt_df = tdt_desc.toPandas()
```

```
[72]: ax = tdt_df.plot(x='loc',y='num_tweets', kind='barh', figsize = (10, 8))


ax.set_ylabel('Location')
```

```

ax.set_xlabel('Total twitter users')
ax.set_title('Top 50 Locations(City/State) of Twitter users of April and May')
ax.invert_yaxis()

# show the plot
plt.show()

```



```
[ ]: # chatgpt and education
```

```

[9]: gpt_key = ['chatgpt', 'gpt', 'ai']
      #filter out rows that do not contain words in keywords
      dt = dt.withColumn('lower', lower(col('text')))
      gpt_dt = dt.filter(col('lower').rlike('|'.join(gpt_key)))

```

```
[81]: gpt_dt.count()
```

```
[81]: 6399921
```

```
[12]: gpt_dt_loc = gpt_dt.filter(col('user_location').isNotNull())
spark.sql("set spark.sql.legacy.timeParserPolicy=LEGACY")
gpt_dt_loc = gpt_dt_loc.select('tweet_id', 'created_at', 'user_location')
gpt_dt_loc = gpt_dt_loc.withColumn("timestamp_unix",
    ↳from_unixtime(unix_timestamp("created_at", "EEE MMM dd HH:mm:ss Z yyyy")))
gpt_dt_loc = gpt_dt_loc.withColumn("date", date_format("timestamp_unix",
    ↳"yyyy-MM-dd"))
gpt_dt_loc = gpt_dt_loc.withColumn('month', month('date'))
```

```
[13]: gpt_ct = gpt_dt_loc.groupby('date').agg(count('*').alias('daily_tweets'))
gpt_ct_desc = gpt_ct.orderBy('date', ascending = True)
```

```
[ ]: gpt_ct_df = gpt_ct_desc.toPandas()
```

```
[102]: gpt_ct_df
```

```
[102]:
```

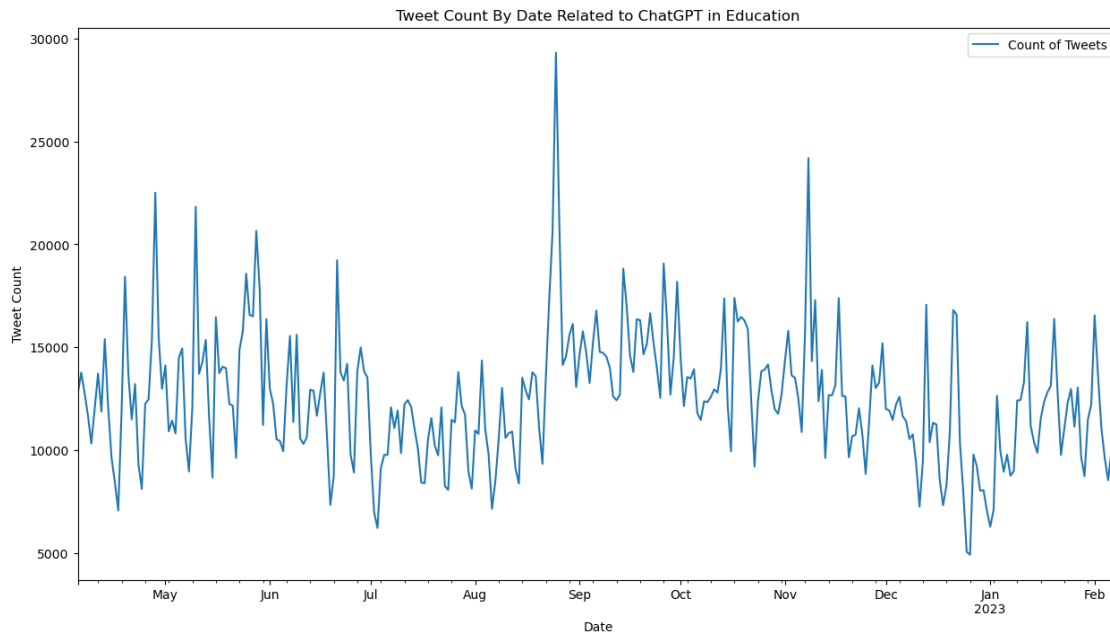
	date	daily_tweets	M
0	2022-04-05	12760	4
1	2022-04-06	13760	4
2	2022-04-07	12778	4
3	2022-04-08	11672	4
4	2022-04-09	10318	4
..
305	2023-02-04	9619	2
306	2023-02-05	8535	2
307	2023-02-06	9963	2
308	2023-02-07	13035	2
309	2023-02-08	9856	2

[310 rows x 3 columns]

```
[103]: ax = gpt_ct_df.plot(kind="line", x='date', y="daily_tweets", figsize=(15, 8))
ax.set_ylabel("Tweet Count")
ax.set_xlabel('Date')
ax.set_title('Tweet Count By Date Related to ChatGPT in Education')

# Adding Legend and Renaming
ax.legend(['Count of Tweets'])

plt.show()
```



```
[96]: gpt_ct_df['date'] = pd.to_datetime(gpt_ct_df['date'])
      gpt_ct_df['M'] = gpt_ct_df['date'].dt.month
      gpt_df_11 = gpt_ct_df[(gpt_ct_df['M'] >= 11) | (gpt_ct_df['M'].isin([1,2]))]
```

```
[97]: gpt_df_11
```

```
[97]:
```

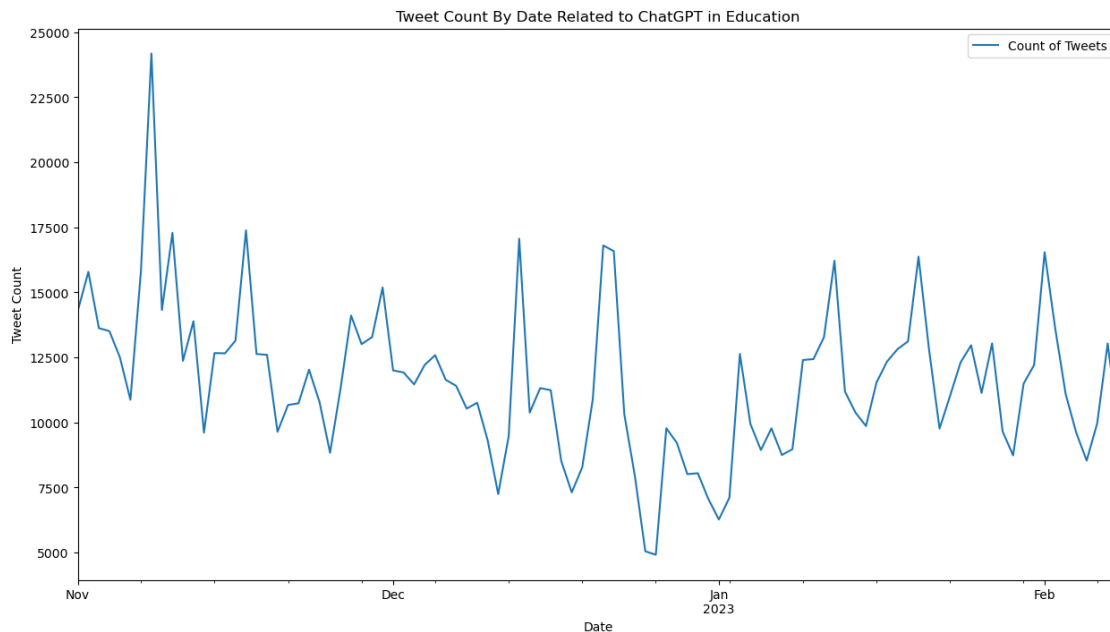
	date	daily_tweets	M
210	2022-11-01	14299	11
211	2022-11-02	15791	11
212	2022-11-03	13626	11
213	2022-11-04	13512	11
214	2022-11-05	12504	11
..
305	2023-02-04	9619	2
306	2023-02-05	8535	2
307	2023-02-06	9963	2
308	2023-02-07	13035	2
309	2023-02-08	9856	2

```
[100 rows x 3 columns]
```

```
[98]: ax = gpt_df_11.plot(kind="line", x='date', y="daily_tweets", figsize=(15, 8))
      ax.set_ylabel("Tweet Count")
      ax.set_xlabel('Date')
      ax.set_title('Tweet Count By Date Related to ChatGPT in Education')
```

```
# Adding Legend and Renaming
ax.legend(['Count of Tweets'])

plt.show()
```



```
[14]: gpt_loc = gpt_dt_loc.filter(col('month')==11)
gpt_loc = gpt_loc.withColumn('loc', split(gpt_loc.user_location, ", ").
    ↳getItem(0))
gpt_loc = gpt_loc.filter(col("user_location").like("%,%"))
gpt_loc_count = gpt_loc.groupby('loc').agg(count('*').alias('num_tweets'))
gpt_loc_desc = gpt_loc_count.orderBy(col('num_tweets').desc()).limit(50)
```

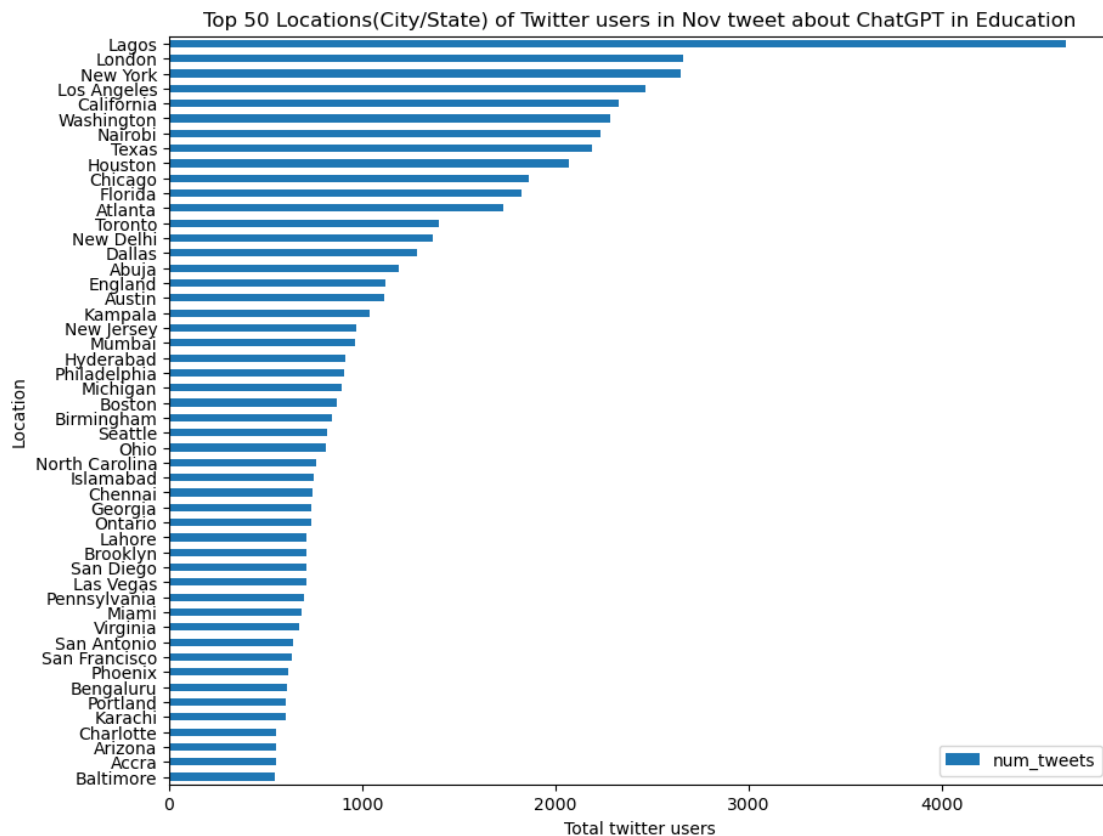
```
[15]: gpt_loc_desc_df = gpt_loc_desc.toPandas()
```

```
[16]: ax = gpt_loc_desc_df.plot(x='loc',y='num_tweets', kind='barh', figsize = (10,
    ↳8))

ax.set_ylabel('Location')
ax.set_xlabel('Total twitter users')
ax.set_title('Top 50 Locations(City/State) of Twitter users in Nov tweet about_
    ↳ChatGPT in Education')
ax.invert_yaxis()

# show the plot
```

```
plt.show()
```



```
[11]: spark.sql("set spark.sql.legacy.timeParserPolicy=LEGACY")
tdt = dt.filter(col('user_location').isNotNull())
tdt = tdt.select('tweet_id', 'created_at', 'user_location')
tdt = tdt.withColumn("timestamp_unix",
    from_unixtime(unix_timestamp("created_at", "EEE MMM dd HH:mm:ss Z yyyy")))
tdt = tdt.withColumn("date", date_format("timestamp_unix", "yyyy-MM-dd"))
tdt = tdt.withColumn('month', month('date'))

tdt_11 = tdt.filter(col('month')==11)
tdt_11 = tdt_11.withColumn('loc', split(tdt_11.user_location, ", ").getItem(0))
tdt_11 = tdt_11.filter(col("user_location").like("%,%"))
tdt_count_11 = tdt_11.groupby('loc').agg(count('*').alias('num_tweets'))
tdt_desc_11 = tdt_count_11.orderBy(col('num_tweets').desc()).limit(50)
tdt_df_11 = tdt_desc_11.toPandas()
ax = tdt_df_11.plot(x='loc',y='num_tweets', kind='barh', figsize = (10, 8))

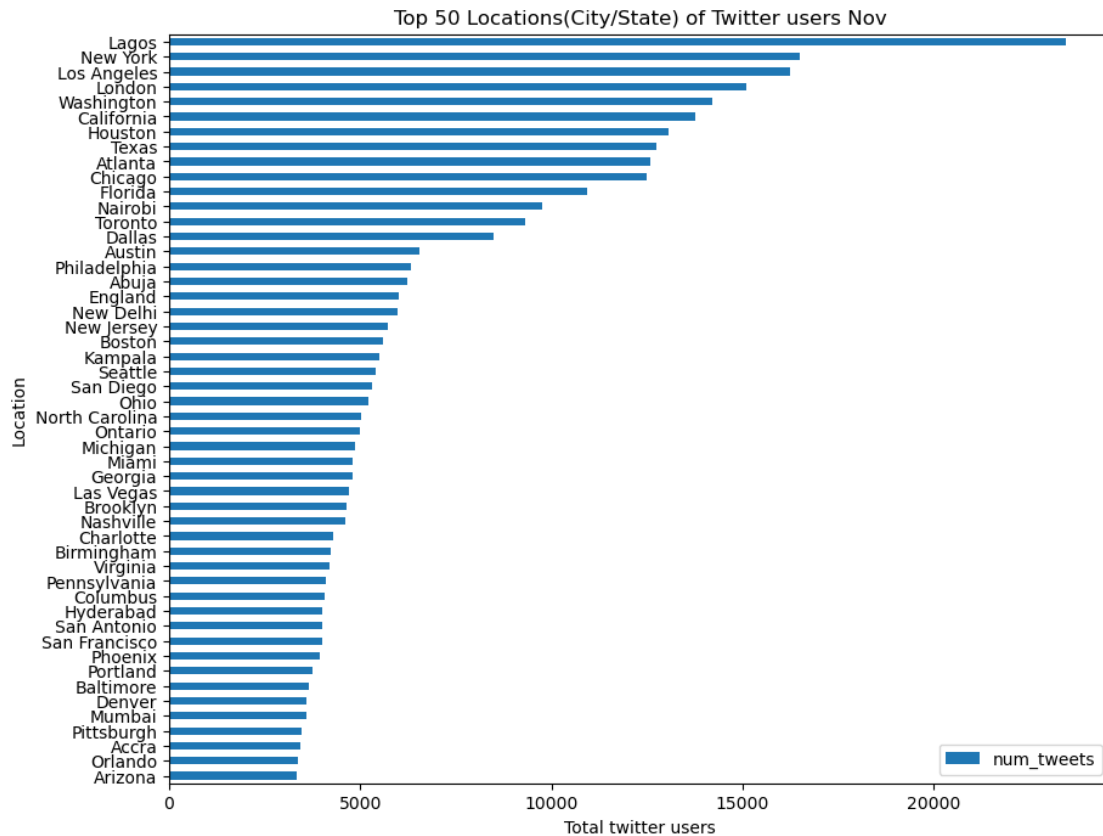
ax.set_ylabel('Location')
ax.set_xlabel('Total twitter users')
```

```

ax.set_title('Top 50 Locations(City/State) of Twitter users Nov')
ax.invert_yaxis()

# show the plot
plt.show()

```



```
[ ]:
```

```

[ ]: gpt_loc_desc.write.format("parquet").\
mode('overwrite').\
save('gs://chen26-bdp/gpt_loc_desc')

```

```
[Stage 59:=====>(1306 + 10) / 1323]
```

```
[108]: t = spark.read.parquet('gs://chen26-bdp/gpt_loc_desc')
```

```
[109]: t
```

```
[109]: +-----+-----+
|      loc|num_tweets|
+-----+-----+
|    Lagos|    10913|
|  New York|    8369|
|   London|    7162|
| Washington|   6949|
| Los Angeles|   6808|
| California|   6516|
|   Nairobi|   6172|
|   Houston|   6007|
|    Texas|   5841|
|   Florida|   5727|
|   Chicago|   5646|
|   Atlanta|   5573|
| New Delhi|   3959|
|    Dallas|   3752|
|   Toronto|   3624|
|   England|   3124|
|    Austin|   2956|
|    Abuja|   2930|
|   Mumbai|   2879|
| New Jersey|   2804|
+-----+-----+
only showing top 20 rows
```

[]:

[]:

[]:

[]:

[]:

[]:

[]:

[]:

[]:

[]:


```
[ ]:
[ ]:
[ ]:
[ ]:
[ ]:
[ ]:
[ ]:
[ ]:
[ ]:
[ ]:
[ ]:
[ ]:
```

```
[81]: %%time
tweets_loc.write.format("parquet").\
mode('overwrite').\
save('gs://chen26-bdp/location_tweets')
```

```
CPU times: user 1.01 s, sys: 335 ms, total: 1.35 s
Wall time: 7min 56s
```

```
[1]: tweets_loc = spark.read.parquet('gs://chen26-bdp/location_tweets')
```

```
[13]: tweet_1 = location.groupby(col('tweet_location')).agg(count('id').
↳ alias('num_tweets'))
```

```
[14]: top50_1 = tweet_1.orderBy(col('num_tweets').desc())
```

```
[ ]:
[ ]:
```


[]: